



```
In [1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [2]: df=pd.read_csv("/content/weatherAUS.csv") # loaded the data set  
df
```

Out[2]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN
...
145455	2017-06-21	Uluru	2.8	23.4	0.0	NaN	NaN
145456	2017-06-22	Uluru	3.6	25.3	0.0	NaN	NaN
145457	2017-06-23	Uluru	5.4	26.9	0.0	NaN	NaN
145458	2017-06-24	Uluru	7.8	27.0	0.0	NaN	NaN
145459	2017-06-25	Uluru	14.9	NaN	0.0	NaN	NaN

145460 rows × 23 columns

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Date              145460 non-null   object  
 1   Location          145460 non-null   object  
 2   MinTemp           143975 non-null   float64 
 3   MaxTemp           144199 non-null   float64 
 4   Rainfall          142199 non-null   float64 
 5   Evaporation       82670  non-null    float64 
 6   Sunshine          75625  non-null    float64 
 7   WindGustDir       135134 non-null   object  
 8   WindGustSpeed     135197 non-null   float64 
 9   WindDir9am        134894 non-null   object  
 10  WindDir3pm        141232 non-null   object  
 11  WindSpeed9am      143693 non-null   float64 
 12  WindSpeed3pm      142398 non-null   float64 
 13  Humidity9am       142806 non-null   float64 
 14  Humidity3pm       140953 non-null   float64 
 15  Pressure9am       130395 non-null   float64 
 16  Pressure3pm       130432 non-null   float64 
 17  Cloud9am          89572  non-null    float64 
 18  Cloud3pm          86102  non-null    float64 
 19  Temp9am           143693 non-null   float64 
 20  Temp3pm           141851 non-null   float64 
 21  RainToday          142199 non-null   object  
 22  RainTomorrow       142193 non-null   object  
dtypes: float64(16), object(7)
memory usage: 25.5+ MB
```

In [5]: `df.describe()`

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine
count	143975.000000	144199.000000	142199.000000	82670.000000	75625.000000
mean	12.194034	23.221348	2.360918	5.468232	7.611178
std	6.398495	7.119049	8.478060	4.193704	3.785483
min	-8.500000	-4.800000	0.000000	0.000000	0.000000
25%	7.600000	17.900000	0.000000	2.600000	4.800000
50%	12.000000	22.600000	0.000000	4.800000	8.400000
75%	16.900000	28.200000	0.800000	7.400000	10.600000
max	33.900000	48.100000	371.000000	145.000000	14.500000

In [6]: `df.isnull().sum()`

Out[6]:

	0
Date	0
Location	0
MinTemp	1485
MaxTemp	1261
Rainfall	3261
Evaporation	62790
Sunshine	69835
WindGustDir	10326
WindGustSpeed	10263
WindDir9am	10566
WindDir3pm	4228
WindSpeed9am	1767
WindSpeed3pm	3062
Humidity9am	2654
Humidity3pm	4507
Pressure9am	15065
Pressure3pm	15028
Cloud9am	55888
Cloud3pm	59358
Temp9am	1767
Temp3pm	3609
RainToday	3261
RainTomorrow	3267

dtype: int64

In [7]: `df.dropna(axis=0, subset=['RainToday', 'RainTomorrow'], inplace=True) # Because`

In [8]: `df.isnull().sum()`

Out[8]:

	0
Date	0
Location	0
MinTemp	468
MaxTemp	307
Rainfall	0
Evaporation	59694
Sunshine	66805
WindGustDir	9163
WindGustSpeed	9105
WindDir9am	9660
WindDir3pm	3670
WindSpeed9am	1055
WindSpeed3pm	2531
Humidity9am	1517
Humidity3pm	3501
Pressure9am	13743
Pressure3pm	13769
Cloud9am	52625
Cloud3pm	56094
Temp9am	656
Temp3pm	2624
RainToday	0
RainTomorrow	0

dtype: int64

In [10]: `df.drop(['Evaporation', 'Sunshine', 'Cloud9am', 'Cloud3pm'], axis=1, inplace=True)`

Why drop them?

Too many guesses needed

Neural networks hate noisy data

Improves model accuracy

```
In [11]: num_cols = df.select_dtypes(include=['float64','int64']).columns  
df[num_cols] = df[num_cols].fillna(df[num_cols].median())
```

```
In [12]: cat_cols = ['WindGustDir','WindDir9am','WindDir3pm']  
  
for col in cat_cols:  
    df[col].fillna(df[col].mode()[0], inplace=True)
```

/tmp/ipython-input-3546851707.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna(df[col].mode()[0], inplace=True)
```

```
In [13]: df['Date'] = pd.to_datetime(df['Date'])  
df['Month'] = df['Date'].dt.month  
df.drop('Date', axis=1, inplace=True)
```

```
In [14]: df.isnull().sum()
```

Out[14]:

	0
Location	0
MinTemp	0
MaxTemp	0
Rainfall	0
WindGustDir	0
WindGustSpeed	0
WindDir9am	0
WindDir3pm	0
WindSpeed9am	0
WindSpeed3pm	0
Humidity9am	0
Humidity3pm	0
Pressure9am	0
Pressure3pm	0
Temp9am	0
Temp3pm	0
RainToday	0
RainTomorrow	0
Month	0

dtype: int64

In [15]:

```
df = pd.get_dummies(  
    df,  
    columns=['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm'],  
    drop_first=True  
)
```

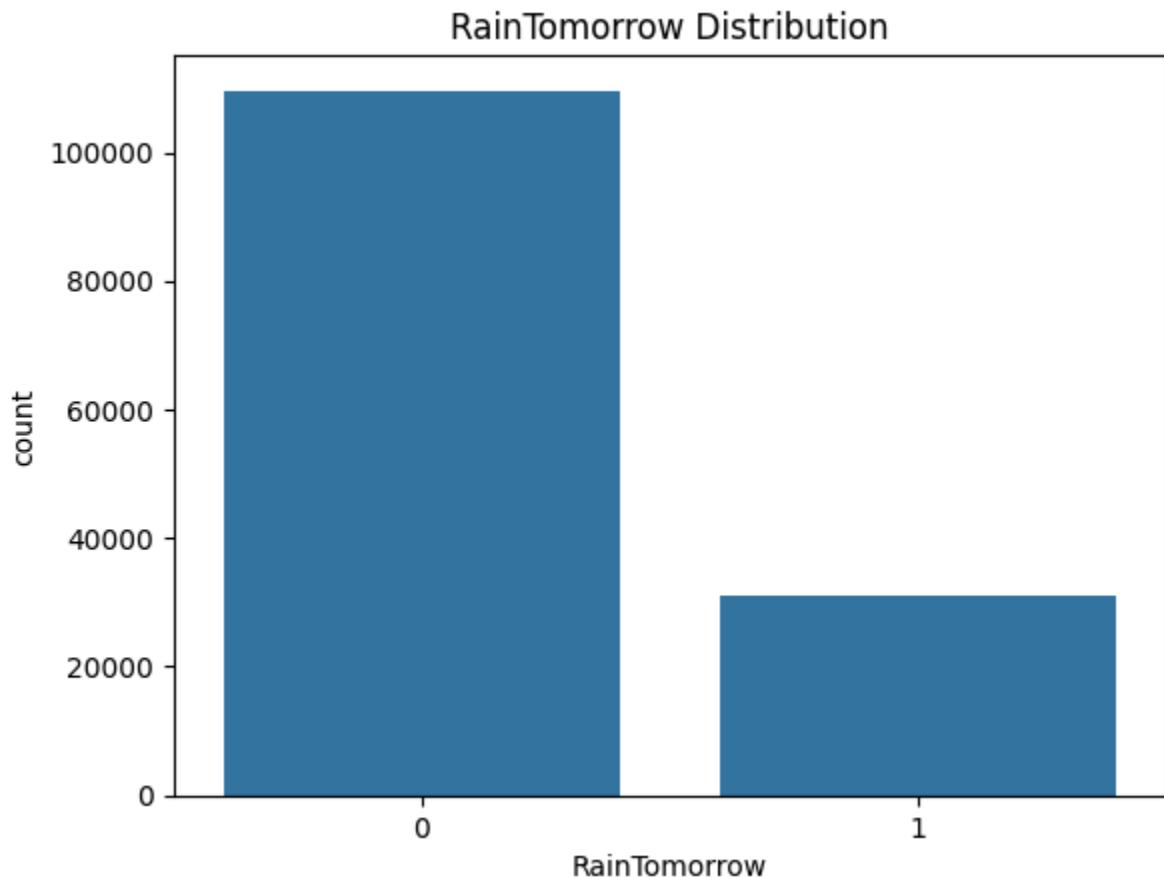
In [16]:

```
from sklearn.preprocessing import LabelEncoder  
  
le = LabelEncoder()  
  
df['RainToday'] = le.fit_transform(df['RainToday'])  
df['RainTomorrow'] = le.fit_transform(df['RainTomorrow'])
```

Target Variable Distribution

```
In [18]: import seaborn as sns
import matplotlib.pyplot as plt

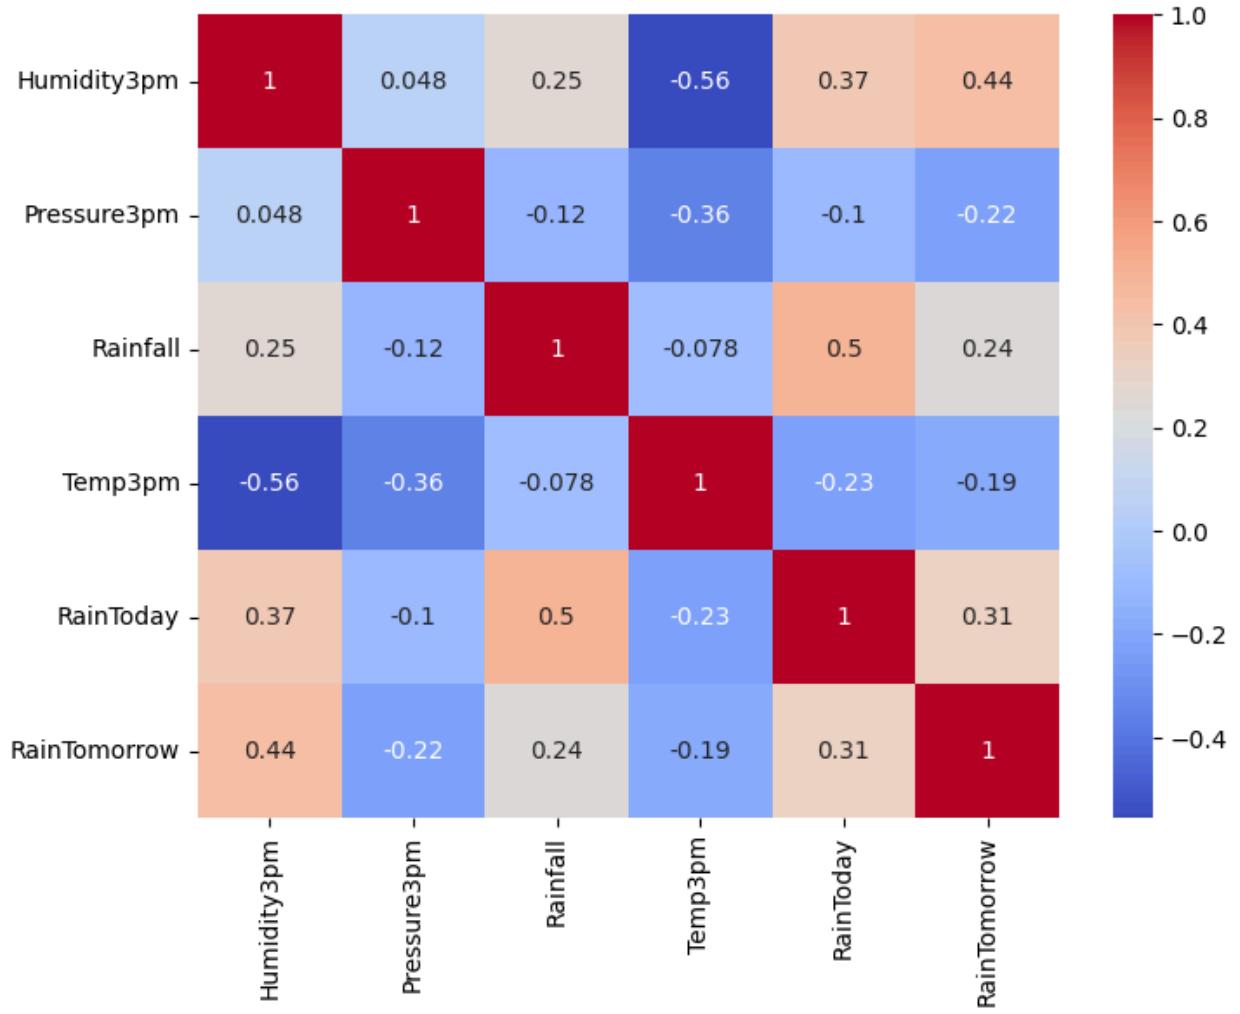
sns.countplot(x='RainTomorrow', data=df)
plt.title("RainTomorrow Distribution")
plt.show()
```



Correlation Heatmap (Feature Relationships)

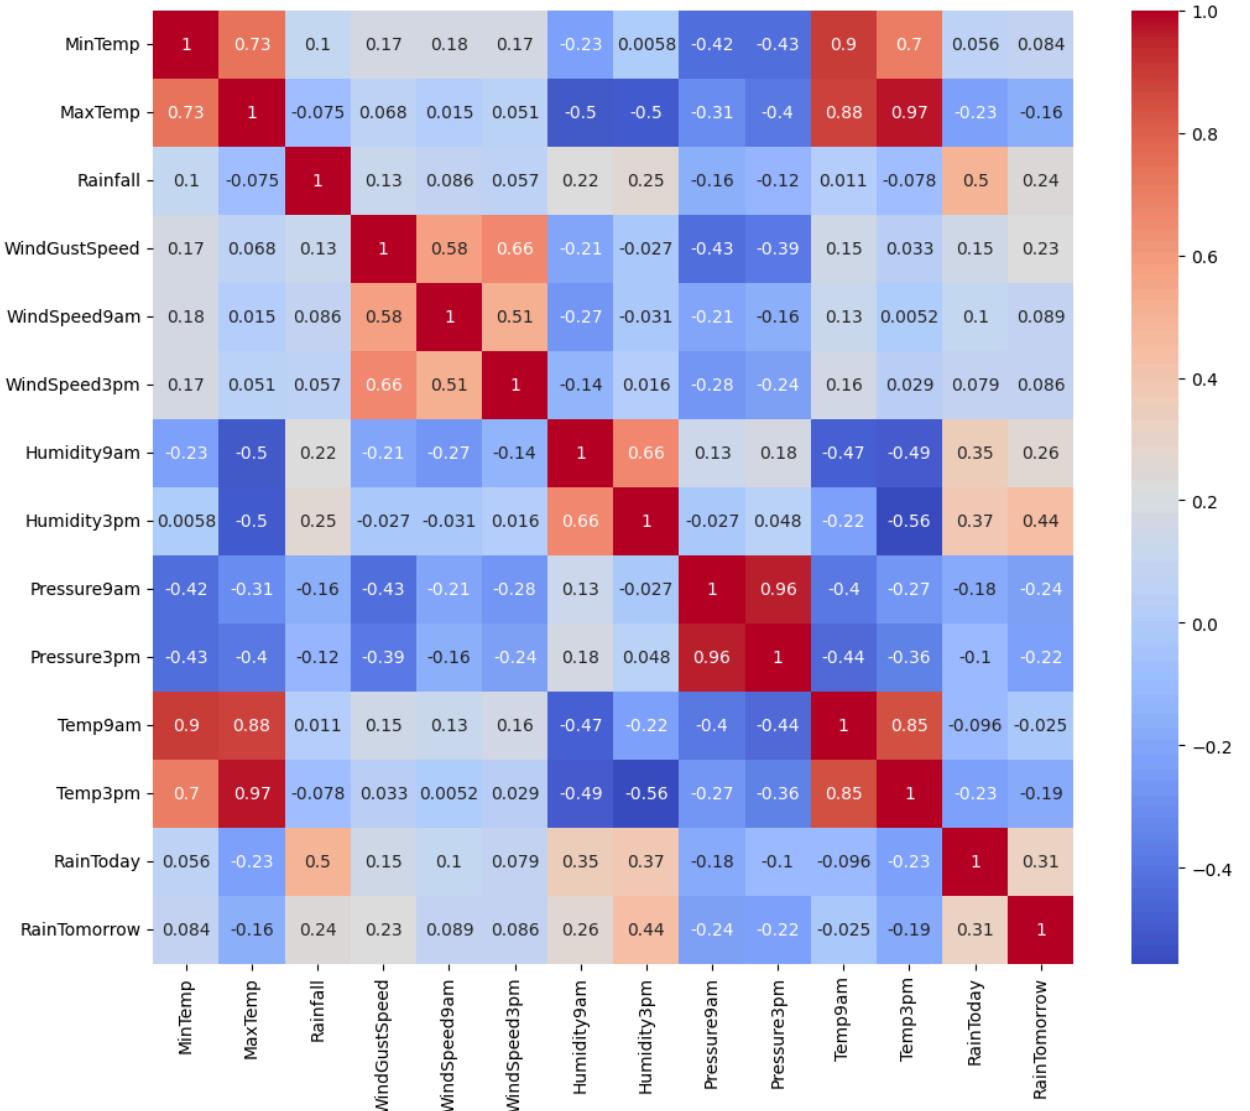
```
In [25]: important_cols = [
    'Humidity3pm', 'Pressure3pm', 'Rainfall',
    'Temp3pm', 'RainToday', 'RainTomorrow'
]

plt.figure(figsize=(8,6))
sns.heatmap(df[important_cols].corr(), annot=True, cmap="coolwarm")
plt.show()
```



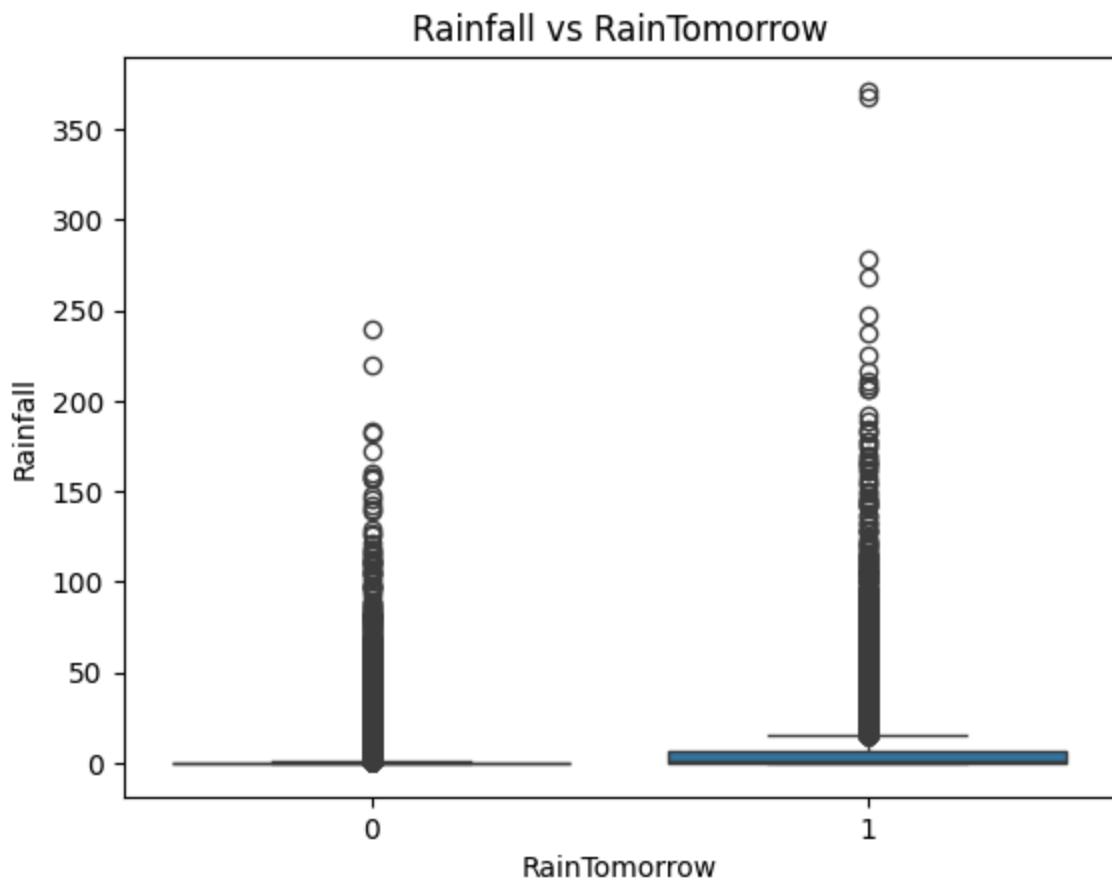
```
In [28]: num_df = df.select_dtypes(include=['float64','int64'])

plt.figure(figsize=(12,10))
sns.heatmap(num_df.corr(), annot=True, cmap="coolwarm")
plt.show()
```



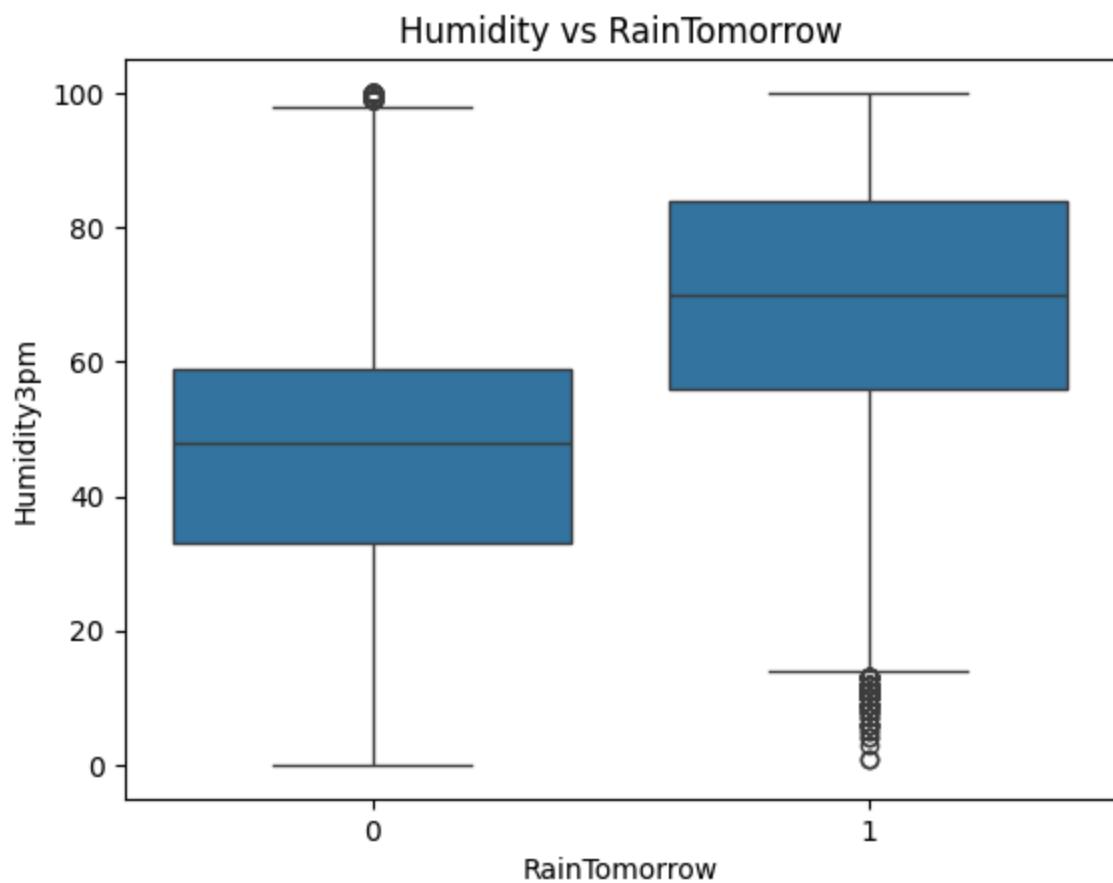
Rainfall vs RainTomorrow

```
In [20]: sns.boxplot(x='RainTomorrow', y='Rainfall', data=df)
plt.title("Rainfall vs RainTomorrow")
plt.show()
```



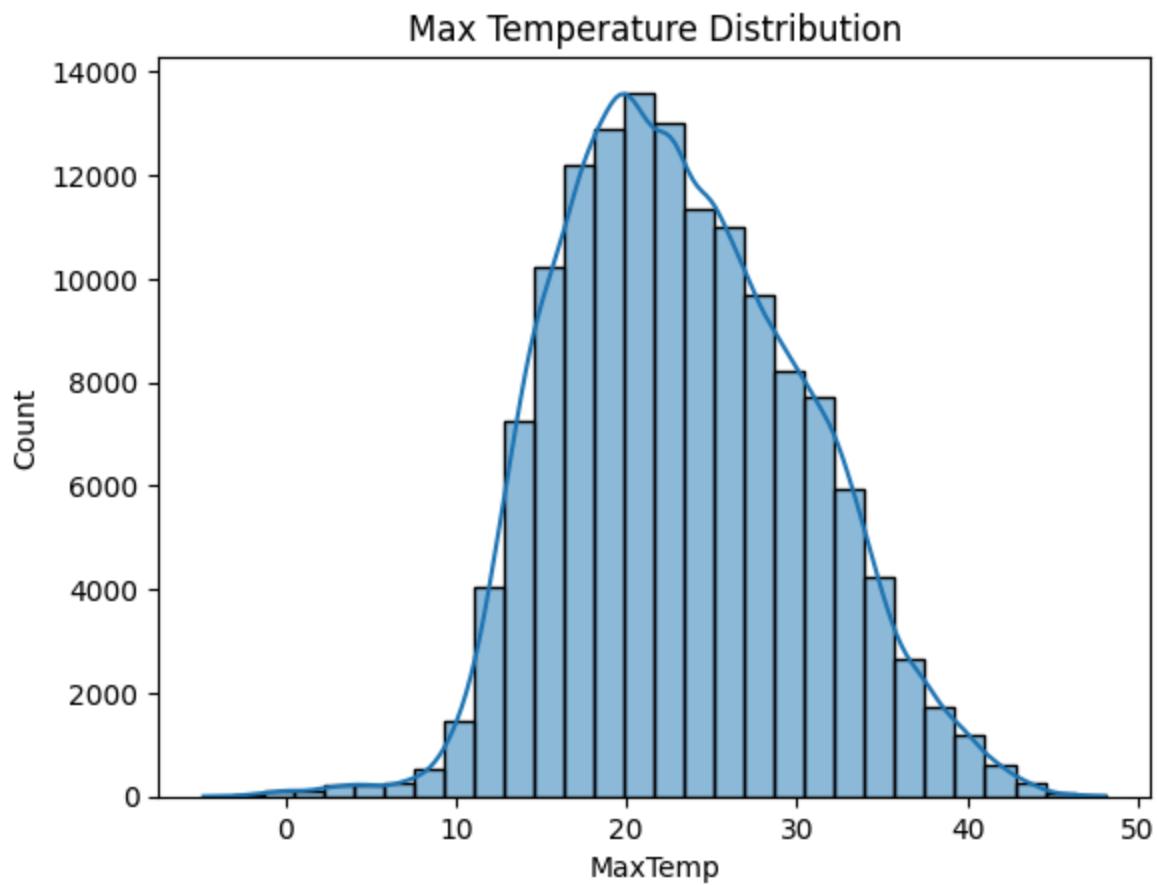
Humidity vs RainTomorrow (Strong Predictor)

```
In [21]: sns.boxplot(x='RainTomorrow', y='Humidity3pm', data=df)
plt.title("Humidity vs RainTomorrow")
plt.show()
```



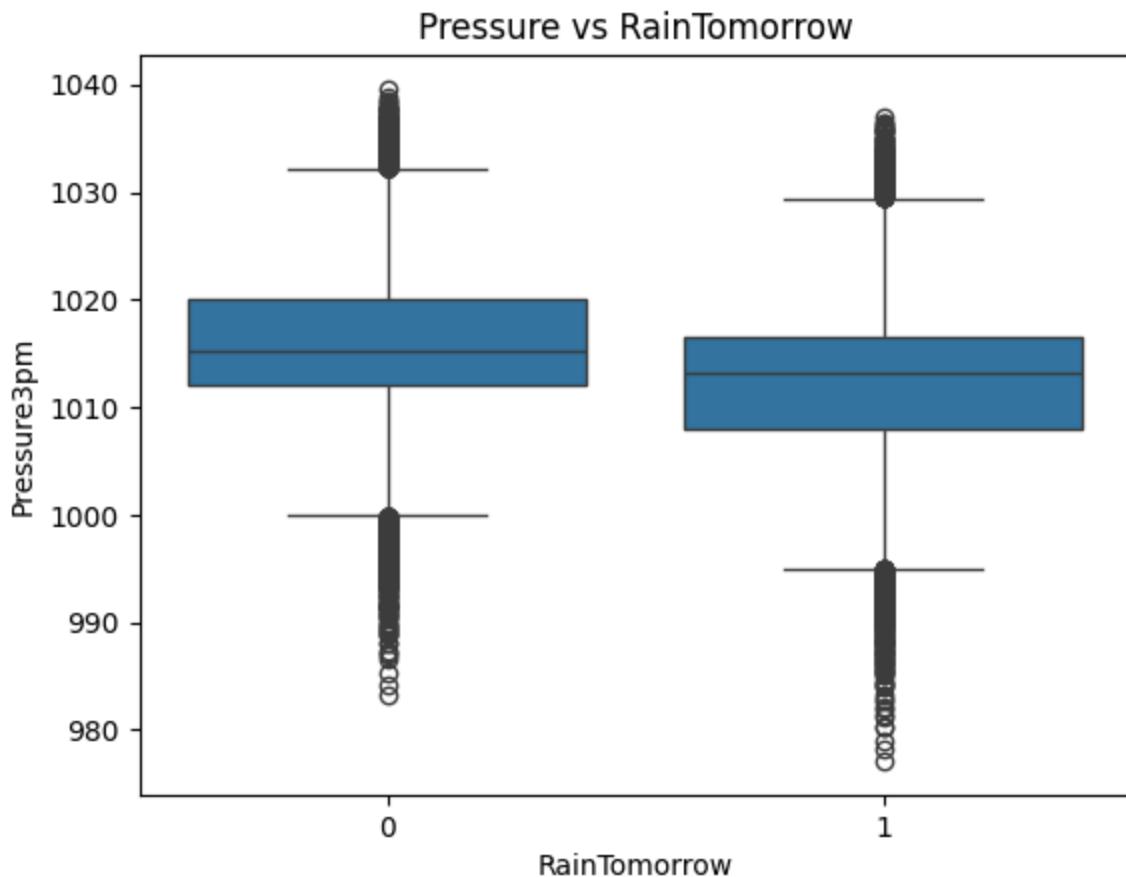
Temperature Distribution

```
In [22]: sns.histplot(df['MaxTemp'], bins=30, kde=True)
plt.title("Max Temperature Distribution")
plt.show()
```



Pressure vs RainTomorrow

```
In [30]: sns.boxplot(x='RainTomorrow', y='Pressure3pm', data=df)
plt.title("Pressure vs RainTomorrow")
plt.show()
```



Split Features and Target

```
In [31]: X = df.drop("RainTomorrow", axis=1)
y = df["RainTomorrow"]
```

```
In [32]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
In [33]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [34]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
model = Sequential()

# Input + Hidden Layer 1
model.add(Dense(64, activation='relu', input_dim=X_train.shape[1]))
```

```
# Hidden Layer 2
model.add(Dense(32, activation='relu'))

# Output Layer
model.add(Dense(1, activation='sigmoid'))
```

```
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/dense.py:93: User
Warning: Do not pass an `input_shape`/`input_dim` argument to a layer. When usi
ng Sequential models, prefer using an `Input(shape)` object as the first layer
in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [35]: model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy']
)
```

```
In [36]: history = model.fit(
    X_train,
    y_train,
    epochs=20,
    batch_size=32,
    validation_split=0.2
)
```

Epoch 1/20
2816/2816 10s 3ms/step - accuracy: 0.8267 - loss: 0.3921 -
val_accuracy: 0.8567 - val_loss: 0.3373
Epoch 2/20
2816/2816 6s 2ms/step - accuracy: 0.8582 - loss: 0.3277 -
val_accuracy: 0.8583 - val_loss: 0.3343
Epoch 3/20
2816/2816 7s 2ms/step - accuracy: 0.8637 - loss: 0.3186 -
val_accuracy: 0.8595 - val_loss: 0.3323
Epoch 4/20
2816/2816 6s 2ms/step - accuracy: 0.8666 - loss: 0.3116 -
val_accuracy: 0.8584 - val_loss: 0.3297
Epoch 5/20
2816/2816 7s 2ms/step - accuracy: 0.8706 - loss: 0.3049 -
val_accuracy: 0.8608 - val_loss: 0.3311
Epoch 6/20
2816/2816 6s 2ms/step - accuracy: 0.8729 - loss: 0.3004 -
val_accuracy: 0.8602 - val_loss: 0.3285
Epoch 7/20
2816/2816 10s 2ms/step - accuracy: 0.8743 - loss: 0.2954 -
val_accuracy: 0.8594 - val_loss: 0.3316
Epoch 8/20
2816/2816 7s 2ms/step - accuracy: 0.8739 - loss: 0.2969 -
val_accuracy: 0.8618 - val_loss: 0.3311
Epoch 9/20
2816/2816 6s 2ms/step - accuracy: 0.8769 - loss: 0.2915 -
val_accuracy: 0.8588 - val_loss: 0.3329
Epoch 10/20
2816/2816 7s 2ms/step - accuracy: 0.8785 - loss: 0.2861 -
val_accuracy: 0.8603 - val_loss: 0.3325
Epoch 11/20
2816/2816 6s 2ms/step - accuracy: 0.8783 - loss: 0.2856 -
val_accuracy: 0.8594 - val_loss: 0.3354
Epoch 12/20
2816/2816 7s 2ms/step - accuracy: 0.8801 - loss: 0.2835 -
val_accuracy: 0.8584 - val_loss: 0.3358
Epoch 13/20
2816/2816 6s 2ms/step - accuracy: 0.8812 - loss: 0.2791 -
val_accuracy: 0.8570 - val_loss: 0.3394
Epoch 14/20
2816/2816 7s 2ms/step - accuracy: 0.8825 - loss: 0.2759 -
val_accuracy: 0.8580 - val_loss: 0.3388
Epoch 15/20
2816/2816 6s 2ms/step - accuracy: 0.8824 - loss: 0.2755 -
val_accuracy: 0.8583 - val_loss: 0.3435
Epoch 16/20
2816/2816 7s 2ms/step - accuracy: 0.8822 - loss: 0.2765 -
val_accuracy: 0.8567 - val_loss: 0.3466
Epoch 17/20
2816/2816 6s 2ms/step - accuracy: 0.8853 - loss: 0.2702 -
val_accuracy: 0.8562 - val_loss: 0.3438
Epoch 18/20
2816/2816 6s 2ms/step - accuracy: 0.8867 - loss: 0.2676 -
val_accuracy: 0.8582 - val_loss: 0.3465

```
Epoch 19/20
2816/2816 10s 2ms/step - accuracy: 0.8873 - loss: 0.2654 -
val_accuracy: 0.8547 - val_loss: 0.3511
Epoch 20/20
2816/2816 7s 2ms/step - accuracy: 0.8885 - loss: 0.2657 -
val_accuracy: 0.8537 - val_loss: 0.3527
```

```
In [37]: loss, accuracy = model.evaluate(X_test, y_test)
print("Test Accuracy:", accuracy)
```

```
880/880 3s 3ms/step - accuracy: 0.8502 - loss: 0.3585
Test Accuracy: 0.8519781231880188
```

```
In [39]: from sklearn.metrics import classification_report
```

```
y_pred_proba = model.predict(X_test)
y_pred = (y_pred_proba > 0.5).astype(int)
print(classification_report(y_test, y_pred))
```

```
880/880 2s 2ms/step
precision    recall   f1-score   support
          0       0.88      0.93      0.91     21897
          1       0.71      0.57      0.63      6261

accuracy                           0.85     28158
macro avg       0.80      0.75      0.77     28158
weighted avg    0.84      0.85      0.85     28158
```

Why the Problem is Classification and Not Regression

In this project, the objective is to predict whether it will rain the next day using the Australian weather dataset. The target variable **RainTomorrow** contains two possible outcomes: Yes or No. Since the output is a discrete category rather than a continuous numerical value, the problem is classified as a **binary classification problem**.

Classification problems involve predicting categorical labels, whereas regression problems involve predicting continuous values such as temperature, rainfall amount, or pressure. In this dataset, the model does not estimate the quantity of rainfall but instead determines the occurrence of rainfall.

Because the target variable is categorical, classification algorithms and evaluation metrics are used. The Artificial Neural Network model uses a sigmoid activation function in the output layer to produce probabilities between 0 and 1, which are then classified into rain or no rain. The model performance is evaluated using classification metrics such as accuracy, confusion matrix, precision, recall, and

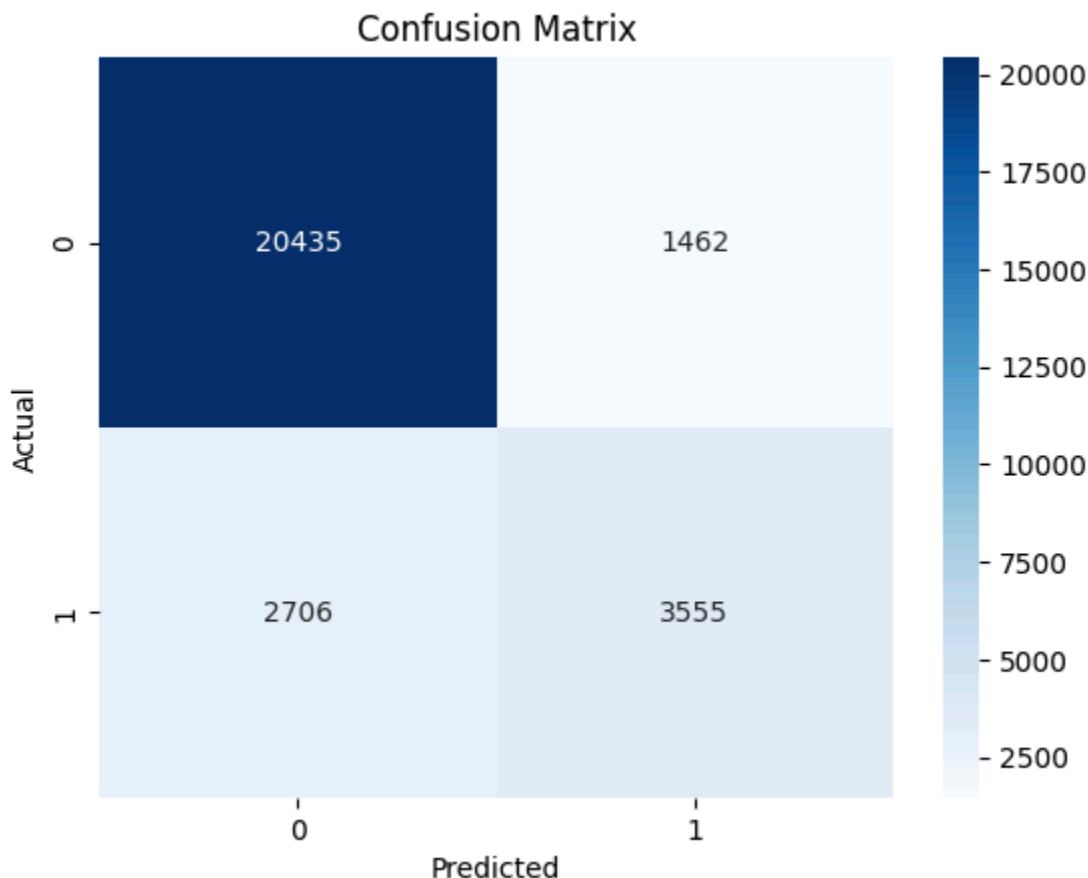
F1-score rather than regression metrics like Mean Squared Error (MSE) or R² score.

Therefore, rainfall prediction in this project is treated as a classification task instead of a regression task.

```
In [41]: import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
In [ ]:
```